

CSC411 Machine Learning

Project 4: Reinforcement Learning

Ariel Kelman

Student No: 1000561368

Gideon Blinick

Student No: 999763000

2 April 2018

1 Introduction

1.1 Results & Report Reproducibility

All results and plots can be generated using the code in `tictactoe.py`. All code is python 3, run with Anaconda 3.6.3. Running the code will save all generated images in the `resources` folder, where they are used by \LaTeX . Note that some of the sections require the code for other sections (in the same file) to be run first. To reproduce the report, simply run it through \LaTeX . This will pull the most recently generated figures from the `resources` folder.

1.2 Tic Tac Toe Environment

The `Environment` class provides the functionality required to play a game of tic tac toe, including against an opponent who plays a random legal move. The game is represented by the `grid` attribute, which is a `numpy` array representing the 9 positions. Each position can have a 0, 1 or 2, representing an empty position, or one filled by player 1 (*X*) or 2 (*O*) respectively. The `turn` attribute can have a value of 1 or 2, and represents which player has the next move. The `done` attribute is a boolean that indicates whether a game has been completed (either with a win, or when the board is full).

The `step()` and `render()` methods allow a tic tac toe game to be played and displayed with text output. Using these methods, a game was played, resulting in the following output:

```
...
.X.
...
=====
O..
.X.
...
=====
O.X
.X.
...
=====
O.X
OX.
```

```

...
=====
O.X
OX.
X..
=====

```

2 Policy

The `Policy` class implements a neural network that learns to play tic tac toe. The provided starter code was modified to be a one-hidden layer neural network; the final code is shown in the code block below.

```

class Policy(nn.Module):
    """
    The Tic-Tac-Toe Policy
    """
    def __init__(self, input_size=27, hidden_size=64, output_size=9):
        super().__init__()
        self.Linear1 = nn.Linear(input_size, hidden_size)
        self.Linear2 = nn.Linear(hidden_size, output_size)

    def forward(self, x):
        h = F.relu(self.Linear1(x))
        out = F.softmax(self.Linear2(h))
        return out

```

In choosing an action, the state is represented as a 27-dimensional vector, using a one-hot encoding type scheme. The first nine elements are 1 if the corresponding location (moving horizontally, and then from top to bottom), and 0 otherwise. Similarly, the next 9 elements are 1 if an *X* (representing player 1) is in the corresponding location; while the last nine elements provide the same functionality for *O* (player 2).

The policy outputs a nine-dimensional vector, which is a probability distribution that is sampled to choose the move for the policy. Thus the policy is stochastic - the `select_action()` function samples this distribution to choose a move.

3 Policy Gradient

The `compute_returns()` function computes the returns based on the reward at the end of a game. The following code block shows how the returns are calculated.

```

def compute_returns(rewards, gamma=1.0):
    """
    Compute returns for each time step, given the rewards
    """
    k = len(rewards)
    rewards = np.array(rewards)
    gammas = np.array([gamma**(i) for i in range(k)])
    G = [sum(rewards[i:]*gammas[:k-i]) for i in range(k)]
    return G

```

The weights are updated on the conclusion of a game...WHY

- 4 Rewards
- 5 Training
- 6 Win Rate
- 7 First Moves
- 8 Limitations